

Global Temperature Anomalies in Practice

An Open, Reproducible Framework for Baseline Harmonisation and ENSO-Aware Visualisation

J. Eduardo Vera-Valdés
Aalborg University, CoRE
eduardo@math.aau.dk

Abstract We present a reproducible pipeline for downloading, processing, and harmonising global temperature anomaly series from HadCRUT5, GISTEMP, NOAAGlobalTemp, Berkeley Earth, and ERA5. All datasets are aligned to a common 1850-1900 pre-industrial baseline to enable direct comparison. The workflow also integrates the Oceanic Niño Index (ONI) to annotate El Niño and La Niña phases in the figures, helping distinguish short-term variability from the long-term warming signal. Final outputs are exported as transparent, analysis-ready CSV files, available through the repository.

1 Introduction

Global temperature anomalies are central to climate variability and change analysis, but practical use across datasets is often limited by inconsistent baseline periods and heterogeneous data formats. Table 1 summarises the baseline periods, formats, and key notes for five of the most widely used global temperature anomaly datasets.

Table 1: Baseline periods and processing considerations for major global temperature anomaly datasets.¹

Dataset	Baseline period	Format	Practical note
HadCRUT5[1]	1961–1990	CSV	Release lag is typically longer than for the other datasets.
NOAAGlobal-Temp[2]	1991–2020	ASCII	Data is presented using space-delimited files. File names are versioned, so the download URL changes over time.
GISTEMP[3]	1951–1980	CSV	Series begins in 1880, so an additional pre-industrial adjustment is needed. Data is in wide format.
Berkeley Earth[4]	1951–1980	TXT	Distributed as space-delimited text files.
ERA5[5]	No fixed baseline	NetCDF	No global average anomaly is provided; users must compute anomalies relative to a selected climatology. Requires setting up a key for accessing the data.

The table highlights two recurring operational barriers: heterogeneous file structures and inconsistent baseline definitions. Together, these barriers reduce comparability, increase processing overhead, and weaken reproducibility in applied climate analysis.

¹Note: baseline definitions may change; always check current provider documentation.

This paper addresses these barriers with a concise, reusable workflow that harmonises all series to a common 1850–1900 reference and a shared tabular format. The implementation is written in Julia [6] for clarity, speed, and straightforward reproducibility.

We standardise all datasets listed in Table 1, compute anomalies relative to the common baseline, and produce a unified analysis table. To support interpretation, we integrate the Oceanic Niño Index [7], [8], [9] and annotate El Niño and La Niña phases in the plots, separating short-term ENSO variability from long-run warming.

The remainder of the paper proceeds from environment setup and source-specific processing to dataset merging and comparative visualisation with ENSO overlays.

2 Necessary Packages and Helper Functions

This notebook uses a Julia project environment with pinned dependencies. The workflow makes use of several packages for data download, processing, and visualisation: `CDSAPI` [10], `CSV` [11], `Dates`, `DataFrames` [12], `HTTP` [13], `NCDatasets` [14], `Plots` [15], and `Statistics`.

`Dates` and `Statistics` are included in the Julia standard library; the remaining packages can be installed with `Pkg.add`. The code snippet below shows how to add the necessary packages. This step only needs to be done once, and the resulting `Project.toml` and `Manifest.toml` files will capture the exact versions used for reproducibility.

```
Pkg.add(["CDSAPI", "CSV", "DataFrames", "HTTP", "NCDatasets", "Plots"])
```

After installation, packages are loaded with `using`.

```
using CDSAPI, CSV, DataFrames, Dates, HTTP, NCDatasets, Plots, Statistics
```

Each dataset follows four steps: (1) download, (2) standardise to date–anomaly format, (3) adjust to 1850–1900 baseline, (4) save as CSV. Helper functions below support this workflow.

A utility function streamlines baseline calculations:

```
function adjust_baseline(df, temp_col; start_year=1850, end_year=1900, date_col=:Date)
    """Compute baseline mean and adjust temperature anomalies."""
    oldbase = mean(df[(df[!, date_col].>=Date(start_year, 1, 1)).&(df[!,
date_col].<Date(end_year, 1, 1)), temp_col])
    return df[!, temp_col] .- oldbase
end
```

3 HadCRUT5

The HadCRUT5 dataset is provided by the Met Office Hadley Centre and the Climatic Research Unit at the University of East Anglia [1]. The HadCRUT5 dataset is available in CSV format from the Met Office website.

```
hfilename = "data/HadCRUT5_global_monthly_average.csv"
open(hfilename, "w") do io
    write(io, HTTP.get("https://www.metoffice.gov.uk/hadobs/hadcrut5/data/HadCRUT.5.1.0.0/
analysis/diagnostics/HadCRUT.5.1.0.0.analysis.summary_series.global.monthly.csv").body)
end
rawhadcrut = CSV.read(hfilename, DataFrame)
rename!(rawhadcrut, :Time => :Date, :Anomaly (deg C)" => :RawTemperature)
```

```

hadcrut = rawhadcrut[!, [:Date, :RawTemperature]]
hadcrut[!, :Temp] = adjust_baseline(hadcrut, :RawTemperature)
CSV.write(hfilename, hadcrut)

```

4 NOAAGlobalTemp

The NOAAGlobalTemp dataset is provided by the National Oceanic and Atmospheric Administration (NOAA) [2]. It is available in ASCII format from the NOAA NCEI website, where space-delimited files are used. The URL of the dataset depends on the month of the data; note the ...YYYYMM.asc name in the URL.

```

nfilename = "data/NOAA_global_monthly_average.csv"
open(nfilename, "w") do io
    write(io, HTTP.get("https://www.ncei.noaa.gov/data/noaa-global-surface-temperature/v6.1/access/timeseries/aravg.mon.land_ocean.90S.90N.v6.1.0.202603.asc").body)
end
# Convert space-delimited ASCII to CSV
write(nfilename, join([join(split(strip(line)), ",") for line in readlines(nfilename)], "\n"))
rawnoaa = CSV.read(nfilename, DataFrame; delim=',', header=0)
dates = Date.(rawnoaa.Column1, rawnoaa.Column2, 1)
noaa = DataFrame(:Date=>dates, :RawTemp=>rawnoaa.Column3)
noaa[!, :Temp] = adjust_baseline(noaa, :RawTemp)
CSV.write(nfilename, noaa)

```

5 GISTEMP

The GISTEMP dataset [3] is available in CSV format from the NASA GISS website. The data is provided in a wide format, where each row corresponds to one year and monthly values are stored in separate columns. To process this data, we need to convert it to a long format, where each row corresponds to a single month. The `longseries` function defined below performs this conversion.

```

function longseries(data)
    """Convert wide-format (year rows) to long-format (monthly rows)."""
    height, last_row = size(data, 1), 12 - count(ismissing, data[end, 2:13])
    many, long = (height - 1) * 12 + last_row, zeros(many, 1)
    for ii in 1:(height-1), jj in 1:12
        long[(ii-1)*12+jj] = data[ii, jj+1]
    end
    for jj in 1:last_row
        long[(height-1)*12+jj] = data[height, jj+1]
    end
    return long
end

```

The function takes a DataFrame in wide format and returns a long-format array. It calculates the total number of months based on the number of years, same as the number of rows, and the number of non-missing months in the last year. It then fills the long-format array by looping through the years and months.

Once the `longseries` function is defined, we can use it to process the GISTEMP data. Note that the data is cleaned from the `***` string used to denote missing values. Furthermore, a date array is created based on the number of months in the dataset, starting from January 1880. The raw temperature anomalies are then adjusted to calculate anomalies relative to the 1850-1900 baseline. Following the data source recommendation, the data are first adjusted to an 1880-1900 baseline and then shifted by 0.038°C to account for pre-1880 conditions².

```
gfilename = "data/GISTEMP_global_monthly_average.csv"
open(gfilename, "w") do io
    write(io, HTTP.get("https://data.giss.nasa.gov/gistemp/tabledata_v4/GLB.Ts%2BdSST.csv").body)
end
longgistemp = CSV.read(gfilename, DataFrame, header=2, missingstring=["***"])
gistemp_long = longseries(longgistemp)[: ]
Tt = length(gistemp_long) - 1
dates = collect(Date(1880, 1, 1):Month(1):Date(1880, 1, 1) + Month(Tt))
gistemp = DataFrame(:Date=>dates, :RawTemp=>gistemp_long)
gistemp[:, :Temp] = adjust_baseline(gistemp, :RawTemp, start_year=1880) .+ 0.038 # +0.038°C
pre-1880 offset per GISS
CSV.write(gfilename, gistemp)
```

6 Berkeley Earth

The Berkeley Earth dataset [4] is available in TXT format from the Berkeley Earth website using space-delimited files.

```
bfilename = "data/BerkeleyEarth_global_monthly_average.csv"
open(bfilename, "w") do io
    write(io, HTTP.get("https://storage.googleapis.com/berkeley-earth-temperature-hr/global/Global_TAVG_monthly.txt").body)
end
rawtemp = CSV.read(bfilename, DataFrame, comment="%", delim=" ", ignorerepeated=true)
rename!(rawtemp, [:Year, :Month, :Anomaly_Monthly, :Unc_Monthly, :Anomaly_Annual, :Unc_Annual, :Anomaly_5yr, :Unc_5yr, :Anomaly_10yr, :Unc_10yr, :Anomaly_20yr, :Unc_20yr])
rawtemp.Date = Date.(rawtemp.Year, rawtemp.Month, 1)
rawtemp.Temp = adjust_baseline(rawtemp, :Anomaly_Monthly)
berkeley = rawtemp[:, [:Date, :Anomaly_Monthly, :Temp]]
rename!(berkeley, :Anomaly_Monthly => :RawTemperature)
CSV.write(bfilename, berkeley)
```

Note that a new variable is created to store the date, which is constructed from the year and month columns in the dataset.

7 ERA5

ERA5 is the fifth generation ECMWF reanalysis for the global climate and weather [5]. Data is available from 1940 in NetCDF format from the Copernicus Climate Data Store.

²<https://data.giss.nasa.gov/gistemp/faq/#q102a>

The API requires setting up a personal key in a `.cdsapirc` file, which should be placed in the root directory of the project. You can get a key by requesting it at the Copernicus Climate Data Store website. The code below assumes that you have already set up the `.cdsapirc` file with your key. To obtain the full dataset, the whole period from 1940 to 2026 is requested, but the file is large, close to 1.5 GB. For update purposes, the code can be modified to request a smaller period, for example, from 2020 to present. The code also specifies the variable to download (2m temperature), the time resolution (monthly), and the data format (NetCDF).

```
# Download ERA5 and compute global mean by month
dataset = "reanalysis-era5-single-levels-monthly-means"
request = ""{"product_type": ["monthly_averaged_reanalysis"], "variable":
["2m_temperature"],
"year": ["1940", "1941", "1942", "1943", "1944", "1945", "1946", "1947", "1948", "1949",
"1950", "1951", "1952", "1953", "1954", "1955", "1956", "1957", "1958", "1959",
"1960", "1961", "1962", "1963", "1964", "1965", "1966", "1967", "1968", "1969",
"1970", "1971", "1972", "1973", "1974", "1975", "1976", "1977", "1978", "1979",
"1980", "1981", "1982", "1983", "1984", "1985", "1986", "1987", "1988", "1989",
"1990", "1991", "1992", "1993", "1994", "1995", "1996", "1997", "1998", "1999",
"2000", "2001", "2002", "2003", "2004", "2005", "2006", "2007", "2008", "2009",
"2010", "2011", "2012", "2013", "2014", "2015", "2016", "2017", "2018", "2019",
"2020", "2021", "2022", "2023", "2024", "2025", "2026"],
"month": ["01", "02", "03", "04", "05", "06", "07", "08", "09", "10", "11", "12"],
"time": ["00:00"], "data_format": "netcdf", "download_format": "unarchived"}""
CDSAPI.retrieve(dataset, request, "data/era5_2m_temperature.nc")
```

The next code snippet computes monthly global mean temperature from gridded ERA5 data. Latitude weights are based on the cosine of latitude to account for grid-cell area differences, so high-latitude cells receive less weight than equatorial cells. ERA5 values are reported in Kelvin and are converted to degrees Celsius.

```
filename = "data/era5_2m_temperature.nc"
ds = NCDataset(filename, "r")
lat_era = ds["latitude"][:]
time_era = ds["valid_time"][:]
date_era = Date.(time_era)

global_mean_C = zeros(length(date_era))
weights = cos.(lat_era * π / 180)
for i in 1:length(date_era)
    t2m = ds["t2m"][:, :, i] # [lon, lat] in K
    global_mean_K = sum(mean(t2m, dims=1)' .* weights) / sum(weights)
    global_mean_C[i] = global_mean_K - 273.15
end
```

Once the monthly global mean temperatures are calculated, anomalies can be computed relative to the 1850-1900 baseline. Unlike the other datasets, ERA5 does not provide a fixed baseline period, so we first compute monthly climatology over 1991-2020 and then derive anomalies relative to that reference. Finally, we add an offset to approximate pre-industrial anomalies and save the processed data.³

³The ERA5 pre-industrial offset (currently 0.88) is obtained from the ERA5 documentation.

```

time_dates = date_era
start_1991 = findfirst(≥(Date(1991,1,1)), time_dates)
end_2020 = findfirst(≥(Date(2021,1,1)), time_dates) - 1
month_clim = [mean(global_mean_C[start_1991:end_2020]
[month.(time_dates[start_1991:end_2020]).==m]) for m in 1:12]
anom_1991_2020 = global_mean_C .- month_clim[month.(time_dates)]
pa_anom = anom_1991_2020 .+ 0.88 # Pre-industrial offset
era5 = DataFrame(Date=time_dates, RawTemperature=anom_1991_2020, Temp=pa_anom)
CSV.write("data/ERA5_global_monthly_average.csv", era5)

```

8 Oceanic Niño Index (ONI)

El Niño (La Niña) is defined as the 3-month running mean of sea surface temperature anomalies in Niño 3.4 region exceeding +0.5°C (or below −0.5°C) for five consecutive overlapping seasons. The data is obtained from the Extended Reconstructed Sea Surface Temperature (ERSST) dataset, which is a global monthly analysis of sea surface temperature data derived from the International Comprehensive Ocean–Atmosphere Dataset (ICOADS) [7], [8]. The ONI data is available in text format from the NOAA Climate Monitoring website [9]. The ONI dataset is provided in a space-delimited format.

```

ofilename = "data/ONI_data.csv"
open(ofilename, "w") do io
    write(io, HTTP.get("https://psl.noaa.gov/data/correlation/oni.data").body)
end
lines = readlines(ofilename);
cleaned_lines = [join(split(strip(line)), ",") for line in lines];
mask = [count(==(','), line) > 11 for line in cleaned_lines];
cleaned_lines = cleaned_lines[mask];
# Convert space-delimited to CSV
write(ofilename, join(cleaned_lines, "\n"));
rawoni = CSV.read(ofilename, DataFrame; delim=',', header=0);
longoni = longseries(rawoni)[:];
longoni = longoni[longoni .!= -99.9]; # Remove missing values;
fechas = Date(1950, 1, 1):Month(1):(Date(1950, 1, 1) + Month(length(longoni)-1));
oni = DataFrame(Date=fechas, Anom=longoni);
CSV.write(ofilename, oni);

```

9 Merge Datasets

Merge all processed datasets on Date to create a unified table for cross-dataset comparison and ENSO-aware trend analysis:

```

hadcrut = CSV.read("data/HadCRUT5_global_monthly_average.csv", DataFrame)
gistemp = CSV.read("data/GISTEMP_global_monthly_average.csv", DataFrame)
noaa = CSV.read("data/NOAA_global_monthly_average.csv", DataFrame)
berkeley = CSV.read("data/BerkeleyEarth_global_monthly_average.csv", DataFrame)
era5 = CSV.read("data/ERA5_global_monthly_average.csv", DataFrame)
oni = CSV.read("data/ONI_data.csv", DataFrame)
# Create date range spanning all datasets

```

```

min_date = minimum([minimum(hadcrut.Date), minimum(gistemp.Date), minimum(noaa.Date),
minimum(berkeley.Date), minimum(oni.Date)])
max_date = maximum([maximum(hadcrut.Date), maximum(gistemp.Date), maximum(noaa.Date),
maximum(berkeley.Date), maximum(oni.Date)])
compiled_data = DataFrame(Date=collect(min_date:Month(1):max_date))

# Merge datasets with standardized column names
for (df, cols, prefix) in [
  (hadcrut, (:RawTemperature, :Temp), "HadCRUT"),
  (gistemp, (:RawTemp, :Temp), "GISTEMP"),
  (noaa, (:RawTemp, :Temp), "NOAA"),
  (berkeley, (:RawTemperature, :Temp), "Berkeley"),
  (era5, (:RawTemperature, :Temp), "ERA5")
]
  df_subset = select(df, :Date, cols[1] => Symbol(prefix*"_RawTemperature"), cols[2]
=> Symbol(prefix*"_Temp"))
  compiled_data = leftjoin(compiled_data, df_subset, on=:Date)
end

# Merge ONI
compiled_data = leftjoin(compiled_data, select(oni, :Date, :Anom => :ONI_Anomaly), on=:Date)
sort!(compiled_data, :Date)
CSV.write("data/Compiled_Global_Temperature_Data.csv", compiled_data)

```

The compiled dataset contains aligned anomalies for all five sources plus ONI, supporting cross-dataset comparison and trend analysis.

10 Plot the Data

Load the compiled dataset and visualize all five temperature series:

```

compiled_data = CSV.read("data/Compiled_Global_Temperature_Data.csv", DataFrame)
theme(:ggplot2)
default(fontfamily = "Computer Modern", tickfontsize = 10, legendfontsize = 10,
titlefontsize = 12)

# Temperature series specifications: (label, column name, marker)
series_specs = [
  ("HadCRUT5", :HadCRUT_Temp, :circle),
  ("GISTEMP", :GISTEMP_Temp, :diamond),
  ("NOAAGlobalTemp", :NOAA_Temp, :+),
  ("Berkeley Earth", :Berkeley_Temp, :xcross),
  ("ERA5", :ERA5_Temp, :star),
]

# Full time series plot
p = plot(
  compiled_data.Date, compiled_data[, series_specs[1][2]],
  title="Global Temperature Anomalies", label=series_specs[1][1],
  xlabel="Date (monthly)", ylabel="Temperature Anomaly (°C)",

```

```

        linewidth=0.5, markersize=1, markershape=series_specs[1][3]
    )

    for (label, col, marker) in series_specs[2:end]
        plot!(compiled_data.Date, compiled_data[:, col], label=label, linewidth=0.5,
              markershape=marker, markersize=1)
    end

    # Add x-axis ticks every 15 years
    plot!(legend=:topleft, xticks=(compiled_data.Date[1:180:end], Dates.format.
    (compiled_data.Date[1:180:end], "Y")))
    display(p)

```




Figure 1: Global Temperature Anomalies

10.1 Zoom into Recent Trends

The plot shows accelerating warming and recent exceedance of the 1.5°C threshold in monthly anomalies. Zoom into the recent 30-year period to better visualize trends and ENSO variability:

```

compiled_zoomed = compiled_data[compiled_data.Date .>= Date(1995, 1, 1), :]

# 30-year plot
p_zoomed = plot(
    compiled_zoomed.Date, compiled_zoomed[:, series_specs[1][2]],
    title="Global Temperature Anomalies (Last 30 Years)", label=series_specs[1][1],
    xlabel="Date (monthly)", ylabel="Temperature Anomaly (°C)",

```

```

        linewidth=0.5, markershape=series_specs[1][3], markersize=1,
    )

    for (label, col, marker) in series_specs[2:end]
        plot!(compiled_zoomed.Date, compiled_zoomed[:, col], label=label, linewidth=0.5,
              markershape=marker, markersize=1)
    end

    plot!(legend=:topleft,          xticks=(compiled_zoomed.Date[1:60:end],          Dates.format.
(compiled_zoomed.Date[1:60:end], "Y")))
    display(p_zoomed)

```




Figure 2: Global Temperature Anomalies (Last 30 Years)

10.2 Adding ENSO Phases

Overlay El Niño and La Niña events on the 30-year plot to distinguish ENSO variability from underlying warming. Classify ONI values into three categories: El Niño ($\geq +0.5^\circ\text{C}$, red shading), La Niña ($\leq -0.5^\circ\text{C}$, blue shading), and Neutral (between).

```

# Classify ONI into El Niño, La Niña, and Neutral phases
ONI_Anom = compiled_zoomed.ONI_Anomaly[.!ismissing.(compiled_zoomed.ONI_Anomaly)]
indicator = zeros{Int, length(OEI_Anom)}
for i in 1:length(OEI_Anom)
    if ONI_Anom[i] >= 0.5
        indicator[i] = 1 # El Niño
    elseif ONI_Anom[i] <= -0.5

```

```

        indicator[i] = -1 # La Niña
    end
end

# Add shaded regions for ENSO phases to the zoomed plot
p_oni = copy(p_zoomed)
i = 1
while i <= length(indicator)
    if indicator[i] in (-1, 1)
        start_idx = i
        while i <= length(indicator) && indicator[i] == indicator[start_idx]
            i += 1
        end
        stop_idx = i - 1
        if stop_idx - start_idx >= 3 # Minimum duration filter
            color = indicator[start_idx] == 1 ? :red : :blue
            vspan!(p_oni, [compiled_zoomed.Date[start_idx], compiled_zoomed.Date[stop_idx]],
                color=color, alpha=0.1, label="")
        end
    else
        i += 1
    end
end
display(p_oni)

```

Global Temperature Anomalies (Last 30 Years)




Figure 3: Global Temperature Anomalies with ENSO Phases

11 Conclusions

This paper provides a reproducible workflow for harmonising major global temperature anomaly datasets to a common 1850-1900 baseline and unified structure. The resulting dataset supports transparent cross-product comparison and consistent downstream analysis.

By combining harmonised anomaly series with ONI-based ENSO annotations, the framework separates short-term variability from the long-term warming signal. The approach is designed for reuse in climate services and policy-facing analysis, including tracking progress against Paris Agreement-related temperature targets. All code and processed datasets are openly published, supporting open-science practices and enabling independent verification and reuse across the research community. The annotated multi-dataset figures are designed for clear visual communication of temperature trends to both scientific and broader audiences.

12 Citation

If you use any of the data or code in this notebook, please cite the original datasets and this notebook as follows:

12.1 Text

Vera-Valdés, J.E. (2026). “Global Temperature Anomalies in Practice”. EarthArXiv (13310) DOI: 10.31223/X56N3Z.

12.2 BibTex

```
@article{vera-valdesGlobalTemperatureAnomalies,  
  title = {Global Temperature Anomalies in Practice},  
  author = {Vera-Valdés, J Eduardo},  
  year = {2026},  
  journal = {EarthArXiv},  
  doi = {10.31223/X56N3Z},  
  url = {https://doi.org/10.31223/X56N3Z}  
}
```

Bibliography

- [1] C. P. Morice *et al.*, “An updated assessment of near-surface temperature change from 1850: The HadCRUT5 data set,” *Journal of Geophysical Research: Atmospheres*, vol. 126, no. 3, p. e2019JD032361, 2021.
- [2] B. Huang, X. Yin, M. J. Menne, R. S. Vose, and H.-M. Zhang, “NOAA Global Surface Temperature Dataset (NOAAGlobalTemp).” [Online]. Available: <https://doi.org/10.25921/rzxg-p717>
- [3] GISTEMP, “GISS Surface Temperature Analysis (GISTEMP), version 4.” [Online]. Available: <https://data.giss.nasa.gov/gistemp/>
- [4] R. A. Rohde and Z. Hausfather, “The Berkeley Earth land/ocean temperature record,” *Earth System Science Data*, vol. 12, no. 4, pp. 3469–3479, 2020, doi: 10.5194/essd-12-3469-2020.
- [5] H. Hersbach *et al.*, “ERA5 monthly averaged data on single levels from 1940 to present.” [Online]. Available: <https://doi.org/10.24381/cds.f17050d7>
- [6] J. Bezanson, A. Edelman, S. Karpinski, and V. B. Shah, “Julia: A fresh approach to numerical computing,” *SIAM review*, vol. 59, no. 1, pp. 65–98, 2017, [Online]. Available: <https://doi.org/10.1137/141000671>

- [7] B. Huang *et al.*, “Extended Reconstructed Sea Surface Temperature, Version 6 (ERSSTv6). Part I: An Artificial Neural Network Approach,” *Journal of Climate*, vol. 38, no. 4, pp. 1105–1121, 2025, doi: 10.1175/JCLI-D-23-0707.1.
- [8] B. Huang *et al.*, “Extended Reconstructed Sea Surface Temperature, Version 6 (ERSSTv6). Part II: Upgrades on Quality Control and Large-Scale Filter,” *Journal of Climate*, vol. 38, no. 4, pp. 1123–1136, 2025, doi: 10.1175/JCLI-D-24-0185.1.
- [9] NOAA National Centers for Environmental Information, “Oceanic Nino Index (ONI).” [Online]. Available: <https://www.ncei.noaa.gov/access/monitoring/enso/>
- [10] M. Y. Chan and J. contributors, “CDSAPI.jl.” [Online]. Available: <https://github.com/JuliaClimate/CDSAPI.jl>
- [11] J. Quinn and J. contributors, “CSV.jl.” [Online]. Available: <https://github.com/JuliaData/CSV.jl>
- [12] M. Bouchet-Valat and B. Kamiński, “DataFrames.jl: Flexible and Fast Tabular Data in Julia,” *Journal of Statistical Software*, vol. 107, no. 4, pp. 1–32, 2023, doi: 10.18637/jss.v107.i04.
- [13] JuliaWeb contributors, “HTTP.jl.” [Online]. Available: <https://github.com/JuliaWeb/HTTP.jl>
- [14] A. Barth, “NCDatasets.jl: A Julia Package for Manipulating netCDF Data Sets,” *Journal of Open Source Software*, vol. 9, no. 101, p. 6504, 2024, doi: 10.21105/joss.06504.
- [15] S. Christ, D. Schwabeneder, C. Rackauckas, M. K. Borregaard, and T. Breloff, “Plots.jl – a User Extendable Plotting API for the Julia Programming Language,” *Journal of Open Research Software*, 2023, doi: 10.5334/jors.431.